

Getting Started with ReplicaSets

This lesson will introduce ReplicaSets to us and we will create a cluster for getting started with ReplicaSets.

WE'LL COVER THE FOLLOWING ^

- Understanding ReplicaSets
- Creating A Cluster

Understanding ReplicaSets

Most applications should be scalable and all must be fault tolerant. Pods do not provide those features, **ReplicaSets** do.

We learned that Pods are the smallest unit in Kubernetes. We also learned that Pods are **not fault tolerant**. If a Pod is destroyed, Kubernetes will do nothing to remedy the problem. That is if Pods are created without **Controllers**.

The first Controller we'll explore is called *ReplicaSet*. Its primary, and pretty much only function, is to ensure that a specified number of replicas of a Pod matches the actual state (almost) all the time. That means that ReplicaSets make Pods scalable.

We can think of ReplicaSets as a *self-healing* mechanism. As long as elementary conditions are met (e.g., enough memory and CPU), Pods associated with a ReplicaSet are guaranteed to run. They provide fault-tolerance and high availability.

If you're familiar with Replication Controllers, it is worth mentioning that ReplicaSet is the next-generation *ReplicationController*. The only significant difference is that ReplicaSet has extended support for selectors. Everything else is the same. ReplicationController is considered deprecated, so we'll focus only on ReplicaSet.

ReplicaSet's primary function is to ensure that the specified number of replicas of a service are (almost) always running.

Let's explore ReplicaSet through examples and see how it works and what it does.

The first step is to create a Kubernetes cluster.

All the commands from this chapter are available in the [04-rs.sh](#) Gist

Creating A Cluster

We'll continue using Minikube to simulate a cluster locally.

```
minikube start --vm-driver=virtualbox  
kubectl config current-context
```



We created a single-node cluster and configured `kubectl` to use it.

Before we explore the first ReplicaSet example, we'll enter into the local copy of the vfarcic/[k8s-spec](#) repository and pull the latest version.

```
cd k8s-specs  
git pull
```



Now that the cluster is running and the repository with the specs is up-to-date, we can create our first ReplicaSet.

In the next lesson, we will create our first ReplicaSet.