

Defining Resource Defaults and Limitations within a Namespace

In this lesson, we will find out why resource defaults and limitations are needed and how to define these.

WE'LL COVER THE FOLLOWING



- Why we Need Resource Defaults and Limitations?
- Defining Default Requests and Limits
 - Creating a Namespace
 - Looking into the Definition
 - `default` and `defaultRequest`
 - `min` and `max`
 - Creating the Resource
 - Looking into the Description
 - Looking into Updated Definition
- Creating Resources
- Looking into the Description

Why we Need Resource Defaults and Limitations?

We already learned how to leverage Kubernetes Namespaces to create clusters within a cluster. When combined with RBAC, we can create Namespaces and give users permissions to use them without exposing the whole cluster. Still, one thing is missing.

We can, let's say, create a `test` Namespace and allow users to create objects without permitting them to access other Namespaces. Even though that is better than allowing everyone full access to the cluster, such a strategy would not prevent people from bringing the whole cluster down or affecting the

performance of applications running in other Namespaces. The piece of the puzzle we're missing is resource control on the Namespace level.

We already discussed that every container should have resource **limits** and **requests** defined. That information helps Kubernetes schedule Pods more efficiently. It also provides it with the information it can use to decide whether a Pod should be evicted or restarted.

Still, the fact that we can specify **resources** does not mean that we are forced to define them. We should have the ability to set default **resources** that will be applied when we forget to specify them explicitly.

Even if we define default **resources**, we also need a way to set limits. Otherwise, everyone with permissions to deploy a Pod can potentially run an application that requests more resources than we're willing to give.

Defining Default Requests and Limits

Our next task is to define default requests and limits as well as to specify minimum and maximum values someone can define for a Pod.

Creating a Namespace

We'll start by creating a test Namespace.

```
kubectl create namespace test
```



Looking into the Definition

With a playground Namespace created, we can take a look at a new definition.

```
cat res/limit-range.yml
```



The **output** is as follows.

```
apiVersion: v1
kind: LimitRange
metadata:
  name: limit-range
spec:
  limits:
  - default:
      memory: 50Mi
      cpu: 0.2
```



```
defaultRequest:
  memory: 30Mi
  cpu: 0.05

max:
  memory: 80Mi
  cpu: 0.5
min:
  memory: 10Mi
  cpu: 0.01
type: Container
```

We specified that the resource should be of `LimitRange` kind. Its `spec` has four `limits`.

`default` and `defaultRequest` #

The `default` limit and `defaultRequest` entries will be applied to the containers that do not specify resources. If a container does not have memory or CPU limits, it'll be assigned the values set in the `LimitRange`. The `default` entries are used as limits, and the `defaultRequest` entries are used as requests.

`min` and `max` #

When a container does have the resources defined, they will be evaluated against `LimitRange` thresholds specified as `max` and `min`. If a container does not meet the criteria, the Pod that hosts the containers will not be created.

Creating the Resource

We'll see a practical implementation of the four `limits` soon. For now, the next step is to create the `limit-range` resource.

```
kubectl --namespace test create \
  -f res/limit-range.yml \
  --save-config --record
```

We created the `LimitRange` resource.

Looking into the Description

Let's describe the `test` Namespace where the resource was created.

```
kubectl describe namespace test
```

The **output**, limited to the relevant parts, is as follows.



```
...
Resource Limits
Type      Resource Min  Max  Default Request Default Limit Max Limit/Request Ratio
-----
Container cpu      10m  500m 50m          200m          -
Container memory   10Mi 80Mi 30Mi          50Mi          -
```

We can see that the `test` Namespace has the resource limits we specified. We set four out of five possible values.

The `maxLimitRequestRatio` is missing and we'll describe it only briefly. When `MaxLimitRequestRatio` is set, container request and limit resources must both be non-zero, and the limit divided by the request must be less than or equal to the enumerated value.

Looking into Updated Definition

Let's take a look at yet another variation of the `go-demo` definition.

```
cat res/go-demo-2-no-res.yml
```



The only thing to note is that none of the containers have any resources defined.

Creating Resources

Next, we'll create the objects defined in the `go-demo-2-no-res.yml` file.

```
kubectl --namespace test create \
  -f res/go-demo-2-no-res.yml \
  --save-config --record

kubectl --namespace test \
  rollout status \
  deployment go-demo-2-api
```



We created the objects inside the `test` Namespace and waited until the deployment `"go-demo-2-api"` was `successfully rolled out`.

Looking into the Description

Let's describe one of the Pods we created.



```
kubect1 --namespace test describe \  
pod go-demo-2-db
```

The **output**, limited to the relevant parts, is as follows.



```
...  
Containers:  
  db:  
    ...  
    Limits:  
      cpu:      200m  
      memory:   50Mi  
    Requests:  
      cpu:      50m  
      memory:   30Mi  
...
```

Even though we did not specify the resources of the `db` container inside the `go-demo-2-db` Pod, the resources are set. The `db` container was assigned the `default` limits of the `test` Namespace as the container limit. Similarly, the `defaultRequest` limits were used as container requests.

As we can see, any attempt to create Pods hosting containers without resources will result in the Namespace limits applied.

We should still define container resources instead of relying on Namespace default limits. They are, after all, only a fallback in case someone forgot to define resources.

In the next lesson, we will explore what happens when the resources are defined but they exceed the Namespace's set limits.