

Observe Metrics Server Data

In this lesson, we will observe the data contained in Metrics Server.

WE'LL COVER THE FOLLOWING

- Memory consumption of pods running in `kube-system`
- Get resource usage of pods using `--all-namespaces`
- Get resource usage of pods using `--containers`
- Flow of data using `kubectl top`
- Scrape the metrics using `JSON`
- `Metrics Server` for machines

Resource usage of the nodes is useful but is not what we're looking for. In this chapter, we're focused on auto-scaling Pods. But, before we get there, we should observe how much memory each of our Pods is using. We'll start with those running in the `kube-system` Namespace.

Memory consumption of pods running in `kube-system`

Execute the following from your command line to see the memory consumption of all the pods running in the Kube-system.

```
kubectl -n kube-system top pod
```

The **output** (on Docker For Desktop) is as follows.

NAME	CPU(cores)	MEMORY(bytes)
etcd-docker-for-desktop	16m	74Mi
kube-apiserver-docker-for-desktop	33m	427Mi
kube-controller-manager-docker-for-desktop	44m	63Mi
kube-dns-86f4d74b45-c47nh	1m	39Mi
kube-proxy-n56hd	2m	22Mi

kube-proxy-r56kd	2m	22Mi
kube-scheduler-docker-for-desktop	13m	23Mi

We can see resource usage (CPU and memory) for each of the Pods currently running in `kube-system`. If we do not find better tools, we could use that information to adjust the `requests` of those Pods to be more accurate. However, there are better ways to get that info, so we'll skip adjustments for now. Instead, let's try to get current resource usage of all the Pods, no matter the Namespace.

Get resource usage of pods using `--all-namespaces`

```
kubectl top pods --all-namespaces
```

The **output** (on Docker For Desktop) is as follows.

NAMESPACE	NAME	CPU(cores)	MEMORY(bytes)
docker	compose-7447646cf5-wqbwz	0m	11M
docker	compose-api-6fbc44c575-gwhxt	0m	14M
kube-system	etcd-docker-for-desktop	16m	74M
kube-system	kube-apiserver-docker-for-desktop	33m	427M
kube-system	kube-controller-manager-docker-for-desktop	46m	63M
kube-system	kube-dns-86f4d74b45-c47nh	1m	38M
kube-system	kube-proxy-r56kd	3m	22M
kube-system	kube-scheduler-docker-for-desktop	14m	23M
metrics	metrics-server-5d78586d76-pbqj8	0m	10M

That **output** shows the same information as the previous one, only extended to all Namespaces. There should be no need to comment on it.

Get resource usage of pods using `--containers`

Often, metrics of a Pod are not granular enough, and we need to observe the resources of each of the containers that constitute a Pod. All we need to do to get container metrics is to add `--containers` argument.

```
kubectl top pods \
  --all-namespaces \
  --containers
```

The **output** (on Docker For Desktop) is as follows.

NAMESPACE	POD		NAM
		CPU(cores)	MEMORY(bytes)
docker	compose-7447646cf5-wqbwz		compos
		0m	11Mi
docker	compose-api-6fbc44c575-gwhxt		compos
		0m	14Mi
kube-system	etcd-docker-for-desktop		etc
		16m	74Mi
kube-system	kube-apiserver-docker-for-desktop		kube-apiserver
		33m	427Mi
kube-system	kube-controller-manager-docker-for-desktop		kube-controller-man
ager		46m	63Mi
kube-system	kube-dns-86f4d74b45-c47nh		kubedn
		0m	13Mi
kube-system	kube-dns-86f4d74b45-c47nh		dnsmas
		0m	10Mi
kube-system	kube-dns-86f4d74b45-c47nh		sidecar
		1m	14Mi
kube-system	kube-proxy-r56kd		kube-proxy
		3m	22Mi
kube-system	kube-scheduler-docker-for-desktop		kube-scheduler
		14m	23Mi
metrics	metrics-server-5d78586d76-pbqj8		metrics-server
		0m	10Mi

We can see that this time, the output shows each container separately. We can, for example, observe metrics of the `kube-dns-*` Pod separated into three containers (`kubedns`, `dnsmasq`, `sidecar`).

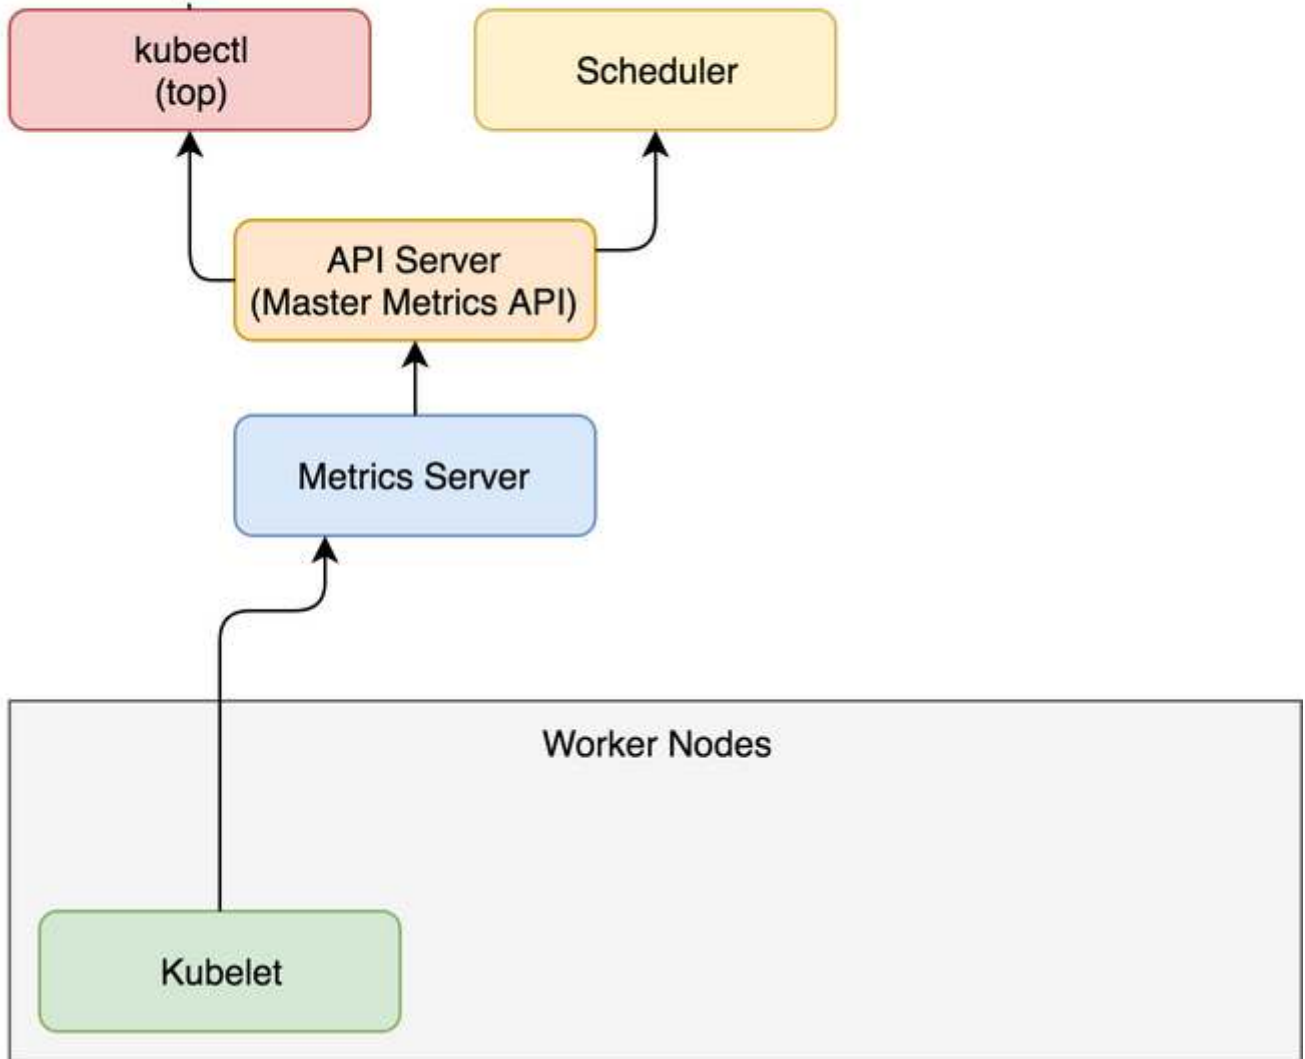
Flow of data using `kubectl top`

When we request metrics through `kubectl top`, the flow of data is almost the same as when the scheduler makes requests. A request is sent to the **API Server (Master Metrics API)**, which gets data from the `Metrics Server`.

server (Master Metrics API), which gets data from the Metrics Server.

which, in turn, was collecting information from Kubelets running on the nodes of the cluster.

"kubectl top" command
retrieves Metrics Server
data through the API



The flow of the data to and from the Metrics Server (arrows show directions of data flow)

Q

`kubectl` command retrieves data from Metrics Server.

Scrape the metrics using **JSON**

While `kubectl top` command is useful to observe current metrics, it is pretty useless if we'd like to access them from other tools. After all, the goal is not for us to sit in front of a terminal with watch `kubectl top pods` command. That would be a waste of our (human) talent. Instead, our goal should be to scrape those metrics from other tools and create alerts and (maybe) dashboards based on both real-time and historical data. For that, we need output in **JSON** or some other machine-parsable format. Luckily, `kubectl` allows us to invoke its API directly in raw format and retrieve the same result as if a tool would query it.

```
kubectl get \  
  --raw "/apis/metrics.k8s.io/v1beta1" \  
  | jq '.'
```

The **output** is as follows.

```
{  
  "kind": "APIResourceList",  
  "apiVersion": "v1",  
  "groupVersion": "metrics.k8s.io/v1beta1",  
  "resources": [  
    {  
      "name": "nodes",  
      "singularName": "",  
      "namespaced": false,  
      "kind": "NodeMetrics",  
      "verbs": [  
        "get",  
        "list"  
      ]  
    },  
    {  
      "name": "pods",  
      "singularName": "",  
      "namespaced": true,  
      "kind": "PodMetrics",  
      "verbs": [  
        "get",  
        "list"  
      ]  
    }  
  ]  
}
```

```
]
}

]
}
```

We can see that the `/apis/metrics.k8s.io/v1beta1` endpoint is an index API that has two resources (`nodes` and `pods`).

Let's take a closer look at the `pods` resource of the metrics API.

```
kubectl get \
  --raw "/apis/metrics.k8s.io/v1beta1/pods" \
  | jq '.'
```

The **output** is too big to display here, so I'll leave it up to you to explore it. You'll notice that the output is `JSON` equivalent of what we observed through the `kubectl top pods --all-namespaces --containers` command.

Metrics Server for machines

There are two important things to note. First of all, it provides current (or short-term) memory and CPU utilization of the containers running inside a cluster. The second and more important note is that we will not use it directly. `Metrics Server` was not designed for humans but for machines. We'll get there later. For now, remember that there is a thing called `Metrics Server` and that you should not use it directly (once you adopt a tool that will scrape its metrics).

Now that we explored `Metrics Server`, we'll try to put it to good use and learn how to auto-scale our Pods based on resource utilization, in the next lesson.