

Switching to Elasticsearch for Storing Metrics

In this lesson, we will see whether we can use Elasticsearch for storing metrics or not.

WE'LL COVER THE FOLLOWING ^

- Using elastic.co to store our metrics
 - Right tool for the job
 - **Prometheus** integration with Kubernetes
 - Using **Prometheus** for logs
 - Using Kibana for metrics

Using **elastic.co** to store our metrics

Now that we had **Elasticsearch** running in our cluster and knowing that it can handle almost any data type, a logical question could be whether we can use it to store our metrics besides logs. If you explore **elastic.co**, you'll see that metrics are indeed something they advertise. If it could replace **Prometheus**, it would undoubtedly be beneficial to have a single tool that can handle not only logs but also metrics. On top of that, we could ditch **Grafana** and keep **Kibana** as a single UI for both data types.

Right tool for the job


Nevertheless, I would strongly advise against using **Elasticsearch** for metrics. It is a general-purpose free-text no-SQL database. That means that it can handle almost any data but, at the same time, it does not excel at any specific format. **Prometheus**, on the other hand, is designed to store time-series data which are the preferred way of exposing metrics. As such, it is more limited in what it does. But, it handles metrics much better than **Elasticsearch**. I believe that using the right tool for the job is better than having a single tool that does too many things, and if you believe the same, **Prometheus is a much better choice for metrics**.

When compared to **Elasticsearch**, and focused only on metrics, **Prometheus** requires much fewer resources (as you already noticed), it is faster, and it has a much better query language. That shouldn't come as a surprise given that both tools are great, but only **Prometheus** is designed to work exclusively with metrics. The increased cost of maintaining an additional tool is well paid off by having a better (and more focused) solution.

Did I mention that notifications generated through **Prometheus** and **Alertmanager** are better than those through **Elasticsearch**?

Prometheus integration with Kubernetes

There's one more important thing to note. **Prometheus** integration with Kubernetes is way better than what **Elasticsearch** offers. That is not a surprise since **Prometheus** is based on the same cloud-native principles as Kubernetes and both belong to [Cloud Native Computing Foundation](#). **Elasticsearch**, on the other hand, comes from a more traditional background.

 *Elasticsearch* is excellent, but it does too much. Its lack of focus makes it inferior to Prometheus for storing and querying metrics, as well as sending alerts based on such data.

Using **Prometheus** for logs

If replacing **Prometheus** with **Elasticsearch** is not a good idea, can we invert the question? Can we use **Prometheus** for logs? The answer is a definite no. As already stated, **Prometheus** is focused only on metrics. If you do adopt it, you need a different tool for storing logs. That can be **Elasticsearch**, **Papertrail**, or any other solution that fits your needs.

Using Kibana for metrics

How about **Kibana**? Can we ditch it in favor of **Grafana**? The answer is yes, but don't do that. While we could create a table in **Grafana** and attach it to **Elasticsearch** as a data source, its capability to display and filter logs is inferior. On the other hand, **Grafana** is much more flexible than **Kibana** for displaying graphs based on metrics. So, the answer is similar to the

Elasticsearch vs. **Prometheus** dilemma. Keep **Grafana** for metrics and use **Kibana** for logs, if you chose to store them in **Elasticsearch**.

Should you add **Elasticsearch** as yet another data source in **Grafana**? If you took previous recommendations, the answer is most likely no. There is not much value in presenting logs as graphs. Even the predefined graph available in **Kibana's Explore** section is, a waste of space. There is no point in showing how many log entries we have in total, nor even how many are error entries. We use metrics for that.

🔍 Logs themselves are too expensive to parse, and most of the time they do not provide enough data to act as metrics.



Elasticsearch is designed to store time-series data which are the preferred way of exposing metrics.

COMPLETED 0%

1 of 1



We saw several tools in action, but we did not yet discuss what we truly need from a centralized logging solution. We'll explore that in the next lesson.