

# EstimatorSpec

Use the EstimatorSpec object to organize training, evaluation, and prediction.

## Chapter Goals:

- Learn how to evaluate a regression model
- Use the `EstimatorSpec` object to organize results from training, evaluation, and prediction

## A. Regression evaluation

Unlike classification models, we can't use the accuracy metric to evaluate regression. Since the output of regression models is a real number, rather than a class prediction, there's no definite way to say what is "correct" or "incorrect".

However, we can tell how good a model output is based on its distance from the corresponding label. For example, if the label for some data observation was 0.2 and the model returned 0.199, the prediction is excellent. On the other hand, if the model returned 812.11, the prediction is likely very poor.

The metric that corresponds to this idea is [mean squared error \(MSE\)](#). The MSE is very similar to the L2-norm, in that both are based on the squared difference between labels and predictions.

In TensorFlow, we obtain the MSE metric using

`tf.metrics.mean_squared_error`. The function takes in the labels and model predictions as its two required arguments.

```
import tensorflow as tf

mse_metric = tf.metrics.mean_squared_error(labels, predictions)
assert isinstance(mse_metric, tuple) and len(mse_metric) == 2
```



Obtaining the MSE metric using `tf.metrics.mean_squared_error`. Note that the function's output is a tuple containing a tensor and an update operation.

The first element of the output tuple is a tensor representing the overall MSE after each evaluation step. The second tuple element is an operation that's used to update the overall MSE after each evaluation step. This gives us a cumulative MSE when we finish evaluating all the data observations.

## B. Using `EstimatorSpec`

The upcoming chapters deal with TensorFlow's `Estimator` API, which encapsulates training, evaluating, and predicting into one compact object. However, in order to use the `Estimator` object, we need to first organize the model results in an `EstimatorSpec`.

The `EstimatorSpec` object is initialized with a single required argument, called `mode`. The `mode` can take one of three values:

- `tf.estimator.ModeKeys.TRAIN`
- `tf.estimator.ModeKeys.EVAL`
- `tf.estimator.ModeKeys.PREDICT`

The keyword arguments required to initialize the `EstimatorSpec` will differ depending on the `mode`.

## Time to Code!

In this chapter you'll be completing the `regressor_fn` by creating helper functions for the `EVAL` and `PREDICT` blocks (the code for the `TRAIN` block is already filled in).

The first helper function you'll create is the `eval_regressor` function, for the `EVAL` block. It applies mean squared error as the evaluation metric.

Then set `mse_metric` equal to `tf.metrics.mean_squared_error` with `labels` as the first argument and `self.predictions` as the second argument.

Set `eval_metric` equal to a dictionary with `'mse'` as the only key and `mse_metric` as the corresponding value.

When initializing `EstimatorSpec` in `EVAL` mode, the required keyword argument is `loss`, meaning that the evaluation will always use the model loss as a metric.

as a metric.

We can add more evaluation metrics through the `eval_metric_ops` keyword argument. The argument takes in a dictionary, which maps string names for each metric to tuple values. Each tuple contains a tensor and update operation (e.g. `mse_metric`).

Set `estimator_spec` equal to `tf.estimator.EstimatorSpec` initialized with `mode` as the required argument. Use `self.loss` for the `loss` keyword argument and `eval_metric` for the `eval_metric_ops` keyword argument.

Then return `estimator_spec`.

```
import numpy as np
import tensorflow as tf

class RegressionModel(object):
    def __init__(self, output_size):
        self.output_size = output_size

    # Helper for regressor_fn
    def eval_regressor(self, mode, labels):
        # CODE HERE
        pass

    # Helper from previous chapter
    def set_predictions_and_loss(self, logits, labels):
        self.predictions = tf.squeeze(logits)
        if labels is not None:
            self.loss = tf.nn.l2_loss(labels - self.predictions)

    # The function for the regression model
    def regressor_fn(self, features, labels, mode, params):
        inputs = tf.feature_column.input_layer(features, params['feature_columns'])
        layer = inputs
        for num_nodes in params['hidden_layers']:
            layer = tf.layers.dense(layer, num_nodes,
                                    activation=tf.nn.relu)
        logits = tf.layers.dense(layer, self.output_size,
                                 name='logits')
        self.set_predictions_and_loss(logits, labels)
        if mode == tf.estimator.ModeKeys.TRAIN:
            self.global_step = tf.train.get_or_create_global_step()
            adam = tf.train.AdamOptimizer()
            self.train_op = adam.minimize(
                self.loss, global_step=self.global_step)
            return tf.estimator.EstimatorSpec(mode,
                                                loss=self.loss, train_op=self.train_op)
        if mode == tf.estimator.ModeKeys.EVAL:
            return self.eval_regressor(mode, labels)
        if mode == tf.estimator.ModeKeys.PREDICT:
            pass
```

Then return `estimator_spec`.

```
logits = tf.layers.dense(layer, self.output_size,
                          name='logits')
self.set_predictions_and_loss(logits, labels)

if mode == tf.estimator.ModeKeys.TRAIN:
    self.global_step = tf.train.get_or_create_global_step()
    adam = tf.train.AdamOptimizer()
    self.train_op = adam.minimize(
        self.loss, global_step=self.global_step)
    return tf.estimator.EstimatorSpec(mode,
        loss=self.loss, train_op=self.train_op)
if mode == tf.estimator.ModeKeys.EVAL:
    # SEE PREVIOUS EXERCISE
    pass
if mode == tf.estimator.ModeKeys.PREDICT:
    return self.predict_regressor(mode, features)
```

