

Training Mode

Chapter Goals:

- Set up the regression function's training code

A. The `EstimatorSpec`

There are three phases to completing the machine learning model: training, evaluation, and prediction. With TensorFlow, we can easily bundle the three phases into a single function using `EstimatorSpec` objects for each phase.

The `EstimatorSpec` object has three modes corresponding to the three phases:

- Training: `tf.estimator.ModeKeys.TRAIN`
- Evaluation: `tf.estimator.ModeKeys.EVAL`
- Prediction: `tf.estimator.ModeKeys.PREDICT`

The `EstimatorSpec` is initialized with whatever mode corresponds to the current phase of completing the machine learning model. For training, the `EstimatorSpec` is also initialized with the model's loss and training operation (for minimizing the loss).

```
inputs = input_layer(features, cols)
```



Creating the model's input layer from a dictionary of feature values (features) and list of feature columns (cols). The `tf.feature_column` prefix is omitted from the `input_layer` function for brevity.

Time to Code!

All code for this chapter goes in the `regression_fn` function.

The `mode` argument for `regression_fn` lets us know whether the model is currently being run for training, evaluation, or prediction. The code for this chapter focuses on model training.

Outside the `if` block from the previous chapter, create another `if` block.

This one should check if `mode` is equal to `tf.estimator.ModeKeys.TRAIN`.

The rest of the code for this chapter goes inside the `if` block.

To keep track of the total number of training steps taken during multiple different training runs, we'll set up a global step object.

Set `global_step` equal to `tf.train.get_or_create_global_step` applied with no arguments.

We'll minimize the model's loss during training using the `ADAM` optimization method, via the `AdamOptimizer` object.

Set `adam` equal to `tf.train.AdamOptimizer` initialized with no arguments.

Set `train_op` equal to `adam.minimize` applied with `loss` as the required argument and `global_step` as the `global_step` keyword argument.

We can now create and return the `EstimatorSpec` object for model training.

Return `tf.estimator.EstimatorSpec` initialized with `mode` as the required argument and `loss` and `train_op` as the `loss` and `train_op` keyword arguments.

```
class SalesModel(object):
    def __init__(self, hidden_layers):
        self.hidden_layers = hidden_layers

    def regression_fn(self, features, labels, mode, params):
        feature_columns = create_feature_columns()
        inputs = tf.feature_column.input_layer(features, feature_columns)
        batch_predictions = self.model_layers(inputs)
        predictions = tf.squeeze(batch_predictions)
        if labels is not None:
            loss = tf.losses.absolute_difference(labels, predictions)
        # CODE HERE

    def model_layers(self, inputs):
        layer = inputs
        for num_nodes in self.hidden_layers:
            layer = tf.layers.dense(layer, num_nodes,
                                    activation=tf.nn.relu)
        batch_predictions = tf.layers.dense(layer, 1)
        return batch_predictions
```



