

Model Training

Chapter Goals:

- Train the regression model

A. Training with the `Estimator`

Since we set up all the code for the regression model, we can now train the model using the `train.tfrecords` file we created in the Data Processing Lab. The `Estimator` object contains a `train` function that lets us easily train the model.

Interestingly, the `train` function's only required argument is a function that takes in no input arguments. This function should set up the input pipeline for the model training.

In our case, it will return the training dataset using the `create_tensorflow_dataset` function from the Data Processing Lab

```
input_fn = lambda:create_tensorflow_dataset('train.tfrecords', 50)
regression_model.train(input_fn)
```

Using the Estimator object (`regression_model`) to run model training. We train with a batch size of 50 and save checkpoints to the `model_ckpt` directory.

Since the `create_tensorflow_dataset` function repeats the dataset indefinitely for training, the above code runs training until we manually stop the process. However, if we want to run training for a fixed number of steps, we can set the `steps` keyword argument.

```
regression_model.train(input_fn, steps=20000)
```

Training for 20000 steps.

Note that the default value for `steps` is `None`, which signifies that training will run until the end of the input dataset. If the input dataset is repeated

indefinitely (as is the case in our training), setting `steps` to `None` will run training until it is manually terminated.

Code for training the regression model is shown below.

```
class SalesModel(object):
    def __init__(self, hidden_layers):
        self.hidden_layers = hidden_layers

    def run_regression_training(self, ckpt_dir, batch_size, num_training_steps=None):
        regression_model = self.create_regression_model(ckpt_dir)
        input_fn = lambda:create_tensorflow_dataset('train.tfrecords', batch_size)
        regression_model.train(input_fn, steps=num_training_steps)

    def create_regression_model(self, ckpt_dir):
        config = tf.estimator.RunConfig(log_step_count_steps=5000)
        regression_model = tf.estimator.Estimator(
            self.regression_fn,
            config=config,
            model_dir=ckpt_dir)
        return regression_model

    def regression_fn(self, features, labels, mode, params):
        feature_columns = create_feature_columns()
        inputs = tf.feature_column.input_layer(features, feature_columns)
        batch_predictions = self.model_layers(inputs)
        predictions = tf.squeeze(batch_predictions)
        if labels is not None:
            loss = tf.losses.absolute_difference(labels, predictions)

        if mode == tf.estimator.ModeKeys.TRAIN:
            global_step = tf.train.get_or_create_global_step()
            adam = tf.train.AdamOptimizer()
            train_op = adam.minimize(
                loss, global_step=global_step)
            return tf.estimator.EstimatorSpec(mode, loss=loss, train_op=train_op)
        if mode == tf.estimator.ModeKeys.EVAL:
            return tf.estimator.EstimatorSpec(mode, loss=loss)
        if mode == tf.estimator.ModeKeys.PREDICT:
            prediction_info = {
                'predictions': batch_predictions
            }
            return tf.estimator.EstimatorSpec(mode, predictions=prediction_info)

    def model_layers(self, inputs):
        layer = inputs
        for num_nodes in self.hidden_layers:
            layer = tf.layers.dense(layer, num_nodes,
                activation=tf.nn.relu)
        batch_predictions = tf.layers.dense(layer, 1)
        return batch_predictions
```

