

# Par/Seq Policy Usage

Let's compare different situations to see where parallel execution works better.

## WE'LL COVER THE FOLLOWING



- Why do you need the sequential policy?
- Limitations and Unsafe Instructions

## Why do you need the sequential policy? #

Most of the time you'll be probably interested in using parallel policy or parallel unsequenced one. But for debugging it might be easier to use `std::execution::seq`. The parameter is also quite convenient as you might easily switch between the execution model using a template parameter. For some algorithms, the sequential policy might also give better performance than the C++14 counterpart.

Read more in the [Benchmark](#) section.

## Limitations and Unsafe Instructions #

The whole point of execution policies is to parallelize standard algorithms in a declarative way effortlessly. Nevertheless, there are some limitations you need to be aware of.

For example with `std::par` if you want to modify a shared resource you need to use some synchronization mechanism to prevent data races and deadlocks [*^readaccess*]:

[*^readaccess*]: When you only want to read a shared resource, then there's no need to synchronise.



```
std::vector<int> vec(1000);
std::iota(vec.begin(), vec.end(), 0);
std::vector<int> output;
std::mutex m;
std::for_each(std::execution::par, vec.begin(), vec.end(),
[&output, &m, &x](int& elem) {
    if (elem % 2 == 0) {
        std::lock_guard guard(m);
        output.push_back(elem);
    }
});
```

Avoid locking inside a parallel operation

The above code filters out the input vector and then puts the elements in the output container.

If you forget about using a mutex (or another form of synchronization), then `push_back` might cause **data races** - as multiple threads might try to add a new element to the vector at the same time.

The above example will also demonstrate weak performance, as using too many synchronization points kills the parallel execution.

When using `par` execution policy try to access the shared resources as little as possible.

With `par_unseq` function invocations might be interleaved, so it's forbidden to use unsafe vectorized code. For example, using mutexes or memory allocation might lead to data races and deadlocks.



```
std::vector<int> vec = GenerateData();
std::mutex m;
int x = 0;
std::for_each(std::execution::par_unseq, vec.begin(), vec.end(),
[&m, &x](int& elem) {
    std::lock_guard guard(m);
    elem = x;
    x++; // increment a shared value
});
```

Unsafe Instructions in Par Unseq Execution

Since the instructions might be interleaved on one thread, you may end up with the following sequence of actions:

```
std::lock_guard guard(m) // for i-th element
std::lock_guard guard(m) // for i+1-th element
...
```

As you can see, two locks (in the same mutex) will happen on a single thread causing a deadlock!

Don't use synchronisation and memory allocation when executing with `par_unseq` policy.

---

In the next lesson, we will take a look at situations in which exceptions arise. Keep on reading to find out more!